

SAM CONSIDINE

Stratford, London | sambconsidine@gmail.com | +44 7808813731 | [LinkedIn](#) | [GitHub](#)

PROFILE

Machine learning engineer with 5 years of industry experience. Passionate about solving challenging real-world problems using technology. Strong background in computer vision, neural search architecture, and language models. Seeking opportunities to increase our collective intelligence by building large-scale AI systems.

SKILLS

Python, PyTorch, Hugging Face (Transformers, Datasets, PEFT, Accelerate, Candle), NLP, PyTorch Geometric, Data Analysis tools (Pandas, Matplotlib, Seaborn, Polars), Big Data (PySpark, AWS Glue, AWS EMR, AWS Athena, Hadoop, HDFS), databases (SQL, PostgreSQL, MongoDB) NumPy, Terraform, Cython, C, C++, Rust, Git.

EDUCATION

University of Cambridge, MPhil Advanced Computer Science

Oct 2021 - Jul 2022

- **Grade:** Distinction — **Thesis:** Neural Algorithmic Reasoning for Pseudotime Trajectory Inference.
- Awarded best Master's thesis of the year by the faculty committee.
- St Edmund's Award for Academic Excellence.

University of York, BSc Mathematics

Sept 2013 - July 2017

- **Grade:** First-Class Honours — **Dissertation:** Computer Vision Based Feature Extraction for Assessing Cell Heterogeneity.
- Winner of the Santander Summer Accelerator Business competition and grant for an idea for an e-commerce platform.

EXPERIENCE

Interim CTO, Noggin HQ

September 2023 – Present

- Designed and built our main credit scoring product in PyTorch: a bidirectional encoder transformer model that uses personal bank transaction records to create neural representations of borrower behaviour.
- Sourced datasets totalling 7 billion bank transaction records to train the model.
- Communication with investors, as well as full evaluation of the business feasibility and tradeoffs of various technical approaches with respect to company objectives.
- Built the company data ETL pipeline to train neural network models at scale using AWS Glue, EMR, Athena and Sagemaker, as well as custom infrastructure on EC2 to perform distributed training.
- Employed 'infrastructure as code' MLOps, using Terraform to deploy an API that could access our credit model through a Sagemaker endpoint.

Machine Learning Engineer, Loci

November 2022 – September 2023

- Built our main product - a multimodal search and recommendation engine for game assets, which was essential in securing our first customers, including one of the world's biggest gaming companies, a crucial component in gaining £4m in seed funding as a team of 4.

- Increased the efficiency of our existing zero-shot game asset classification system by 3 orders of magnitude while improving code simplicity by creating a custom vector store + caching system.
- Re-implemented and fine-tuned large multimodal language/image transformers (OpenCLIP-2B) on 10m novel views of 3D assets, building a data pipeline, allowing us to train, benchmark and deploy models.
- Mentored a team of 6 Masters students in a collaboration with Imperial College to build a system for 3D texture synthesis from a text prompt using stable diffusion, UNet and some custom rendering.

Research Assistant, University of Cambridge

July 2022 – October 2022

- Hired to continue research on my thesis and to build a platform to be used by clinicians to use machine learning systems to assist with diagnosis.
- Worked with some of the top researchers in explainable AI, developing a UI to bridge the knowledge gap between doctors and data scientists. Allowing doctors to use private, explainable AI methods with their patients.

Senior Data Scientist, Arca-Blanca

May 2021 – October 2021

- Worked for a Machine Learning and Data Science consulting firm, consulting with executives to create and assist with predictive modelling tasks throughout their enterprise.
- Lead various projects, collaborating with executive at large companies.
- Lead data scientist of a project for one of the UK's largest software companies, creating a model of customer churn and customer segmentation analysis. The segmentation is still used as a core part of the marketing team strategy.

Machine Learning Engineer, After the Off

September 2017 - January 2021

- Led developer on our main horse racing product: using machine learning on live data feeds to generate in-play odds for horse races, establishing us as a leading provider of in-play horse racing odds.
- Designed and led a project using computer vision to derive horse positions on a race track in real time using only a live video feed of the race. Built the whole pipeline starting from the video feed.

Research Assistant, University of York

July 2017 - September 2017

- Hired to continue work on my dissertation, building computer vision based feature extraction to aid in cancer research.

Self-Employed, Uptodata

June 2012 – December 2016

- Started a successful marketing company as a teenager, providing marketing data to wholesalers.
- Made contacts in the giftware industry after finding a catalog of wholesalers from an event organised by my father, used this to build a database of giftware retailers in the UK, which allowed me to sell marketing information to wholesalers.
- Paid my way through university, allowing me to do without a maintenance loan or other job.